

Aangan Studio — AI Phone Enquiry Agent: Research, Plan & Workflow

Oct 7, 2026 · @Anudeep

Summary

We build a voice agent that picks up Aangan's studio line within seconds, 24/7, qualifies the caller against `qualified.md`, books a consultation, and hands the designer a note with every answer already captured. Existing-client complaints and anything that sounds like a commitment go to a human.

The why is Nikhil's own data: about 96 of ~200 monthly enquiries (48%) get no reply within 48 hours, a third arrive outside 10am to 7pm, and an answer inside one hour converts at 4x a next-day answer. At ₹8-14 lakh per project, a handful of recovered conversions a month pays for the system many times over.

The Cut is Nikhil's third ask: quoting per-square-foot pricing. `pricing.md` is marked "indicative only, varies by site", and a number spoken in the studio's voice is heard as a promise (Check 06: NO). The agent instead explains why a number needs a site visit and books the free consultation. In September, both phone callers who pushed for a ballpark (T02, T13) booked anyway after that exact explanation.

Scope is inbound phone only. WhatsApp and the web form get a written extension path, not a build.

Ownership: what I take responsibility for

I own the outcome, not just the build: every call answered, every qualified caller booked, and a dashboard that proves it. If Nikhil asks "why you?", the answer is that this problem sits between three groups (front desk, designers, Nikhil) and none of them owns the gap between the phone ringing and a designer calling back.

What I own, concretely:

- **The research and decisions:** platform, handoff channel, compliance wording, cost model. Each decision is written down with the reason, so Nikhil can challenge it.
- **The rules the agent follows:** turning `qualified.md` and `services.md` into a script, and getting Nikhil to sign that script off before go-live.

- **The test bar:** the agent must handle all 20 September phone transcripts correctly before it takes a live call.
- **The numbers:** baseline on day one, weekly review, monthly cost report.
- **The off switch:** I can route the line back to the front desk in under a minute if anything goes wrong.

Activity	Nikhil	Me (founder's office)	Designers	Front desk	Agent
Answer calls, any hour	Informed	Accountable	—	Backup in hours	Does it
Qualification rules	Approves	Drafts, maintains	Consulted	Consulted	Follows
Booking consultations	—	Accountable	Own their calendars	Reschedules	Books
Handoff note	—	Designs template	Act on it	—	Writes it
Price and quotes	Owns	—	Give at consultation	—	Never quotes
Existing-client complaints	Owns	Monitors	Responds	Logs	Escalates only
Dashboard and cost	Reads it	Owns	—	—	Logs data

The research I personally still have to close before build (tracked in the Risks section): the telecom route for the existing studio number, Nikhil's sign-off on qualified.md as a script, designer calendar access, and the exact consent wording for recording.

Field research: what the outside evidence changes

The research backs Nikhil's 4x figure and changes three things in the plan: two "free" tools in the stack are not free for a business, outbound AI callbacks are a regulatory risk, and caller data must stay off any tier that trains on it.

Speed-to-lead evidence

Nikhil's 4x is conservative next to published studies. An HBR analysis found firms that contacted a lead within an hour were nearly 7x as likely to qualify it as firms that waited one hour more, and over 60x as likely as those waiting 24 hours or longer. The

MIT/InsideSales Lead Response Management study put contact odds 100x higher at 5 minutes than at 30 minutes ([summary of both studies](#)). Both are US datasets, so I use Nikhil's own 4x for the business case and cite these only as direction.

Voice AI platforms in India

Platform	What it is	Price signal	Role for us
Vaani Labs	Maharashtra-based; real-time voice agents in 40+ languages across phone, web and WhatsApp; Twilio and SIP; an MCP server with a Claude Skill	Free tier, no card; per-minute rate not published	Primary (course stack). Same agent can later serve WhatsApp
Bolna	Mumbai, India-first; Exotel, Plivo and Airtel SIP telephony	About \$0.06/min on pilots; ₹1.5/min platform fee if we bring our own models; Indian numbers \$5/month	Price benchmark and fallback
Vapi / Retell	Global orchestration layers	Vapi \$0.05/min platform fee; Retell \$0.07–0.31/min all-in	Not chosen: USD billing and India last-mile latency

Open question for Vaani: the all-in per-minute rate for Indian inbound calls. I will get this in writing before committing.

Regulation

Rule	What it says	What we do
DPDP Act and Rules 2025	Rules notified Nov 2025; notice and consent duties enforceable 18 months later, May 2027	Build it in now: a one-line spoken notice at call start, data used only for the enquiry, a retention limit, delete on request

Rule	What it says	What we do
TRAI TCCCPR amendment, Feb 2025	Promotional calls from the 140 series, service and transactional calls from the 1600 series; normal 10-digit numbers restricted for telemarketing	The agent only answers inbound calls. No AI-dialled callbacks in v1; a human calls back missed and dropped calls from the studio number
Commercial SMS (DLT registration, from prior knowledge)	Business SMS in India needs a registered header and templates	v1 confirms bookings by calendar invite and email, not SMS
Gemini API terms	Free-tier content is used to improve Google's products; paid-tier content is not	Caller data only on a paid tier

We also disclose the agent at the start of every call ("Aangan's virtual assistant"). That is a trust decision, not a legal claim.

Stack gotchas

Component	Finding	Decision
Vercel Hobby	Restricted to non-commercial personal use; a studio dashboard is commercial	Budget Vercel Pro (\$20/month)
Cal.com / Calendly	Round-robin across designers needs the paid Teams plan, about \$15 per user per month. For 14 designers that is about \$210/month	Use the Google Calendar API (free): read each designer's free/busy, pick by our own rotation rule in Supabase
Resend free	3,000 emails/month, 100/day	Enough for booking confirmations at ~100–200 a month
Gemini Flash-Lite	\$0.30 in / \$2.50 out per million tokens; Gemini 3.8 Flash doubles in price on 1 Jan 2027	Post-call classification costs well under ₹1 a call. Pin the model and watch the January repricing

What the 20 September phone calls tell us

Of 20 phone enquiries, 12 were worth a designer's time, 6 were not a fit, 1 was an existing-client complaint and 1 was missed after hours. The front desk already runs a consistent script; the agent's job is to run that same script at 10:47pm (T08) and never lose the note (T16).

Call	Situation	Class	What the agent must do
T01	3BHK Kothrud, 1,400 sq ft, full redesign, by March; referral	Qualified	Book; record referrer
T02	2BHK Wakad, 950 sq ft, full; asked for a price twice	Qualified	Explain why no number yet; book
T03	Home office in Nashik	Not a fit: area	Decline politely; say why
T04	Ideas for a living room, no execution	Not a fit: advisory only	Explain full-service model; log as future
T05	4BHK Koregaon Park, 2,400 sq ft; friend of Nikhil's network	Qualified, VIP	Book site visit; flag referral to Nikhil
T06	800 sq ft bare-shell office, Baner; founder calling	Qualified, commercial	Book; note headcount and December go-live
T07	Living room and kitchen before Diwali, 3 weeks away	Not a fit: timeline	Explain 8–10 weeks; offer a November start
T08	Missed call at 10:47pm; call-back failed	Missed	Would be answered by the agent
T09	Existing client, no reply from designer for 5 days	Escalate	No qualifying; senior callback within 15 min
T10	1BHK Kharadi, kitchen + bedroom, ₹1–1.5 lakh budget	Not a fit: budget	Never name a floor; route to human review
T11	Rented 2BHK Baner, landlord agrees, no structural work	Qualified	Book; note landlord consent
T12	5,500 sq ft villa, Kalyani Nagar; wants principal designer on site	Qualified, high value	Book principal designer; flag to Nikhil

Call	Situation	Class	What the agent must do
T13	3BHK Aundh, 1,100 sq ft; pushed for a ballpark	Qualified	Explain why no number yet; book
T14	Son calling for parents' new 3BHK in Hadapsar	Qualified, proxy	Book only with owners attending
T15	2BHK Undri, possession in 6 weeks, full home	Qualified	Book site visit
T16	Repeat caller whose Monday enquiry was lost	Qualified, recovery	Recognise number; apologise; book; flag frustration
T17	Call dropped, caller rang back in 2 minutes	Qualified	Merge both calls into one record; book
T18	180 sq ft coworking pod	Not a fit: size	Decline; commercial minimum is 500 sq ft
T19	Restaurant interiors	Not a fit: type	Decline; hospitality is out of scope
T20	2BHK Magarpatta, 900 sq ft, January start; both owners attend	Qualified	Book

Qualification rules inferred from the front desk

These are reverse-engineered from the transcripts. qualified.md is the authority; I reconcile the two and get Nikhil's sign-off before build.

1. **Location:** Pune and PCMC only. Edge areas (Talegaon, in form F03) go to a human.
2. **Service:** design and execution together. No advisory-only visits.
3. **Type:** homes and offices. No restaurants, gyms or other hospitality.
4. **Size:** commercial between 500 and 3,000 sq ft.
5. **Timeline:** design takes 3–4 weeks; one room takes at least 8–10 weeks end to end. A shorter deadline gets offered a later start.
6. **Budget:** if a caller volunteers a budget far below scope, a human reviews it. The agent never states a minimum.
7. **Decision-maker:** the owners take part in the consultation. A relative may call on their behalf.

8. Rentals: fine with landlord consent and reversible work.

The detail that defends the Cut

Both phone callers who pushed for a price (T02, T13) booked once they heard why a number needs a site visit, because they were answered live. The two price-only enquiries that were lost (WhatsApp W02, form F07) got no reply for two days or none at all. Speed closed the gap, not the number.

Automation Brief: the why

Pain. About 96 enquiries a month get no answer for 48 hours, and about 66 arrive when nobody is at the desk. Two people cover three channels for nine hours a day. Notes get lost (T16), and designers spend their first call re-asking what the front desk already knew.

User. The caller first: a Pune homeowner or small-office founder deciding which studio to trust, often comparing three (form F06). Then the designer who receives the handoff, and Nikhil, who reads the numbers.

Outcome. Every call answered in seconds, day or night. Qualified callers leave the call with a booked consultation. The designer gets a complete note within two minutes of hang-up. Nikhil never pulls numbers by hand again.

The journey, today and tomorrow

Step	Today	Tomorrow
1. Call arrives	Answered 10am–7pm only, if a desk person is free. 1 in 3 enquiries lands after hours (T08)	Agent picks up within seconds, any hour, and says it is Aangan's virtual assistant
2. Questions	Desk asks area, scope, size, timeline, owner, in their own words	Agent asks the same questions in a fixed order, in English, Hindi or Marathi
3. Price question	Desk explains why there is no number yet (T02, T13)	Agent uses the same approved explanation; never a number
4. Fit decision	Desk judges on the spot	Agent applies the 8 written rules; unclear cases go to a human queue

Step	Today	Tomorrow
5. Booking	"Let me check the calendar and revert" (T12)	Agent books a designer's slot during the call
6. Record	In someone's head or notebook; sometimes lost (T16)	Structured record in Supabase and a HubSpot deal, automatically
7. Handoff	Forwarded message, no context; designer re-asks	Designer gets a Telegram note with every answer and the recording
8. Measurement	Nikhil pulls numbers by hand once a month	Live dashboard, including what the tool costs

Nine Checks and the Cut

Six checks pass cleanly, two pass only partly, and Check 06 fails for one ask: quoting price. That failure is the Cut. (The brief's task line says "Arjun"; it means Nikhil.)

#	Check	Verdict	Evidence from the case
01	Problem real?	Yes	48% unanswered for 48+ hours; 4x conversion under one hour. Nikhil's own numbers
02	Workflow repeated?	Yes	~200 enquiries a month; the first questions (area, scope, size, timeline, owner) are asked in nearly every September call
03	Input available?	Partly	services.md and qualified.md cover answering and qualifying. pricing.md says "indicative only, varies by site", so it cannot be spoken
04	Output valuable?	Yes	A booked consultation on a ₹8-14 lakh project is directly revenue

#	Check	Verdict	Evidence from the case
05	Impact measurable?	Yes	Time to answer, % answered under one hour, consultations booked, show rate. All can be baselined from September
06	Failure risk OK?	Yes for answering and qualifying; No for pricing	A wrong qualify is recoverable by a designer. A spoken price is heard as a commitment and is not
07	Judgment protected?	Partly	Answering and qualifying can be assisted. Quoting, committing and complaints (T09) stay human
08	ROI worth it?	Yes	One extra project a month at ₹8 lakh is at least 25x the estimated monthly running cost, even at the high estimate (see Cost model)
09	Owner clear?	Yes, once bounded	Agent qualifies and books; the designer owns the relationship; Nikhil owns the price; I own the system

The Cut: no per-square-foot quotes

Nikhil asked for it; Checks 03, 06 and 07 stop it. A number said in the studio's voice before anyone has seen the site is a promise the studio may not keep, and the front desk already knows this: "the difference between a standard and a premium kitchen alone can be 3x" (T13).

What we build instead:

- An approved explanation, in Nikhil's words, of why the number comes at the free consultation.
- If the caller volunteers a budget, it goes into the designer's note, so the first meeting starts there.

- The designer gives the real number at the end of the consultation, as the desk already promises.


When Nikhil pushes back that people won't wait, the answer is in September's calls: answered live, both price-pushers booked (T02, T13).

Components map

The design rule: the AI talks and extracts, plain code decides. The eight qualification rules run as deterministic code on the extracted fields, so a fit decision can be audited and changed without retraining anything.

Case 3 · Aangan phone agent — components map

The AI talks and writes up; plain code decides fit; people own price, complaints and the client.



components map · 5 actors × 6 stages

The two AI boxes sit in the AI column: Vaani's model runs the conversation live, and Gemini writes up the call after hang-up. Every decision (fit, booking, routing) happens in the System lane as plain code, and price appears in no lane at all.


Stage	What it is in this build	Tool
Trigger	An inbound call to the studio number. Pilot: forwarded after hours and on no-answer. Later: every call	Telco call forwarding → Vaani Labs number
Input	The caller's speech, the caller ID, the time of call	Vaani Labs (speech-to-text)
Context	services.md; qualified.md as 8 rules; the approved price explanation; designers' free/busy; past calls from this number; Nikhil's VIP referrer list	Supabase, Google Calendar API
Processing	Live: look up caller, check slots, book, transfer to human. After the call: extract fields to JSON, apply the rules, route	Vaani tool calls; a small API on Vercel
AI	Live conversation model inside Vaani. Post-call extraction and classification by Gemini Flash-Lite. Claude writes the script and rules at design time and builds the code in Claude Code. No AI on price, complaints or final escalation	Vaani, Gemini (paid tier), Claude
Output	Caller: spoken confirmation plus calendar invite. Designer: Telegram note. Nikhil: escalation alert and dashboard. CRM: HubSpot deal for qualified calls. Front desk: queue for unclear cases	Telegram Bot API, HubSpot, Resend, Vercel dashboard

Call workflow

Every call follows one path: answer, recognise, route, qualify, book, hand off. Only three moments need a person: a complaint, an unclear case, and the designer accepting the booking.

Every call ends in one of five places; a person acts in three

Outlined: the agent. Shaded: a person acts. A caller can ask for a person at any point.



inbound call flow · 5 endings, 3 human steps

The two diamonds are the only decisions. The first is the model reading intent; the second is plain code applying the 8 rules to what the caller said.

Call script

1. **Greet and disclose:** "Namaste, Aangan Studio. I'm Aangan's virtual assistant, and this call is recorded so our designers have your details. How can I help?" The agent follows the caller into English, Hindi or Marathi.
2. **Route:** a new project, an existing project, or something else.
3. **Ask, in this order:** where the home or office is; what they want done; rough carpet area; current state (bare, lived-in, awaiting possession); timeline; who will decide and attend; how they heard of Aangan.
4. **If price comes up:** the approved explanation. A figure before a designer sees the site would mislead, because materials alone can triple the cost of one kitchen; the consultation is free and ends with a real number.
5. **Read back and book:** "So that's a 3BHK in Kothrud, about 1,400 sq ft, a full redesign done by March. I have Thursday at 11 or Saturday at 10:30 with a designer. Which suits you?"
6. **Close:** confirm name and email for the invite, and tell the caller the designer will already know everything they shared.

Designer handoff: Telegram, with HubSpot as the record

I choose Telegram for the designer's note and keep email for the caller's confirmation. The whole system exists to cut response time from days to minutes; a note that sits in an inbox undoes that at the last step.

Factor	Telegram bot	Email (Resend)
Speed to the designer	Push notification on the phone they already carry	Lands in an inbox among everything else
Two-way	Accept / Can't take it buttons on the message	No reliable read or accept signal
Measurable	Time from booking to designer accept goes straight to the dashboard	Not measurable without extra tooling
Formal record	Weak on its own	Strong
Cost	Free	Free up to 3,000/month, 100/day

The accept button matters most. T09 is a designer who went silent for five days. With an accept button, a note not accepted within 30 minutes during working hours moves to the next designer and alerts the design lead, so silence is caught on day one, not day five.

The two real objections, and the answers:

- **"Our designers use WhatsApp, not Telegram."** True for most Indian teams. Telegram is a one-time install for 14 people; WhatsApp's business API needs Meta approval and template setup, which belongs with the WhatsApp channel build later.
- **"Client data on a personal chat app."** The note carries only what the designer needs to prepare. HubSpot holds the full record, and the recording link opens only behind a login.

Handoff note template

NEW CONSULTATION · Qualified · High value
 Anand Sharma · Kalyani Nagar · Villa, G+2, ~5,500 sq ft
 Booked: Fri 18 Sep, 11:00am · Site visit · Principal designer asked for
 Scope: full end-to-end design; new construction, empty
 Timeline: move in by March; "no rush"
 Decision-maker: caller is the owner
 Budget: not discussed · Asked about price: no
 Flags: high value, Nikhil notified
 Call: 7 min 30 s · recording and transcript in HubSpot
 [Accept] [Can't take it]

The designer opens the consultation already knowing every answer from the call, which is Nikhil's second ask met in full.

Dashboard: what Nikhil sees

One page on Vercel that answers three questions without anyone pulling numbers: are calls being answered, are they turning into consultations, and what does it cost per booked consultation. Every number comes from the Supabase call log, HubSpot and logged usage, never from manual entry except fixed monthly fees.

Metric	Definition	September baseline	Target
Answered under 1 hour	Share of phone enquiries with a live answer within 60 min	About 52% of all enquiries answered within 48 hours; phone-only figure to pull from telco call records	100%
Time to answer	Median seconds from ring to pickup	Not tracked	Under 10 seconds
Qualified rate	Qualified calls / all enquiry calls	12 of 20 (60%) in the September sample	Track, no target
Booked on the call	Qualified calls with a slot booked before hang-up	Several were "I'll revert" (T05, T12)	80% or more
Designer accept time	Booking to Accept tap, working hours	Not tracked	Under 30 min
Re-asked questions	Designer's yes/no after each consultation	Common, per Nikhil	Under 10% "yes"
Show rate	Consultations attended / booked	Unknown	Track
Escalations closed	Complaints with a senior callback inside 15 min	T09 handled by hand	100%
Cost per booked consultation	All tool costs this month / consultations booked	n/a	Track
Pipeline from agent calls	HubSpot deal value from agent-sourced calls	n/a	Track

The cost panel shows this month's spend by line (voice minutes, numbers, AI tokens, hosting), cost per call and cost per booked consultation, so the tool reports its own price as Nikhil asked. Voice minutes come from Vaani's usage records; AI tokens are logged per call by our own code.

A quality panel shows 10 random calls a week for me to review, each with the agent's decision and whether I would overturn it. The overturn rate is the agent's error rate.

Cost model

The system should cost about ₹11,000 a month at mid estimates, and under ₹30,000 even at the most expensive voice rate. Voice minutes are almost the whole bill; everything else is free tier or a flat fee.

Assumptions: 200 calls a month (the slide's figure; the September sample suggests about half of 200 enquiries are phone, so this is a ceiling), 5 minutes average (September's answered calls averaged 4.9 minutes), and ₹88 to the dollar.

Line	Basis	Low (₹/month)	Mid (₹/month)	High (₹/month)
Voice AI, all-in	1,000 min at \$0.06 / \$0.10 / \$0.31 per min	5,280	8,800	27,280
Phone number	About \$5 a month	440	440	440
Post-call AI	Gemini Flash-Lite, ~6,000 tokens in and 600 out per call	60	60	60
Dashboard hosting	Vercel Pro, \$20	1,760	1,760	1,760
Supabase, HubSpot, Google Calendar, Telegram, Resend, GitHub	Free tiers	0	0	0
Total		7,540	11,060	29,540

At the mid estimate that is about ₹55 a call. If roughly 96 consultations are booked a month (60% qualified, 80% of those booked), it is about ₹115 per booked consultation.


One extra won project a month at ₹8 lakh covers the high estimate 27 times over. Two costs sit outside the table: forwarded calls may be billed by the studio's telco as an outgoing leg (to check on Aangan's plan), and my own time, about 4 hours a week for call review and rule updates.

Build plan

The agent can be answering every call from 12 November if Nikhil signs the rules this week. No phase starts until the gate before it is passed; a failed gate adds a week, not a shortcut.

Three gates stand between today and an agent on every call

Build plan, October to November 2026. Dates assume Nikhil signs the rules this week.



build plan · 4 phases, 3 gates

The pilot spans Diwali, when enquiry volume may dip; if fewer than about 30 calls reach the agent, the pilot runs a third week before Gate 3.

Claude Code sessions

- Foundation:** repo on GitHub; Supabase tables for calls, callers, bookings, designers and costs.
- Agent:** the Vaani agent with the signed script and four tools: look up caller, check slots, book, transfer to a person.
- Booking:** Google Calendar free/busy across designers, our own rotation rule, and the calendar invite by email.
- After the call:** webhook → Gemini extracts fields to a fixed JSON schema → the 8 rules run as code → routing.
- Handoff:** Telegram bot with Accept and Can't-take-it buttons and the 30-minute reassign; HubSpot deal sync.
- Dashboard:** the Vercel page, with every call's token and minute cost logged.
- Test harness:** replay the 20 September phone calls as scripted test calls, plus the 10 hard cases below; fix and repeat until all pass.

The rules engine is also run on all 40 September enquiries, including WhatsApp and form, because their fields test the rules even though only the phone goes live.

The 10 hard cases for Gate 2

Case	Based on	Pass means
Price push, English	T02	Explains, books, says no number
Price push, Hinglish	W03 style	Same, in Hinglish
Angry existing client	T09	Transfers or promises a senior callback; no questions asked
Son calling for parents	T14	Books with owners attending
Repeat caller, lost note	T16	Recognises the number and apologises
Call drops mid-way	T17	Second call continues the first record
Edge location	F03 Talegaon	Goes to the human queue
Budget far below scope	T10	No floor stated; human review
Restaurant or gym	T19, W07	Declines kindly, says why
"Is this a robot? I want a person."	New	Confirms, transfers or books a callback

Risks, guardrails and open questions

The worst failure is the agent saying a number or mishandling an angry client; both have a hard guardrail and a weekly check.

Risk	Guardrail
Agent states a price or a commitment	Hard rule in the script; a post-call scan flags any rupee amount or "per sq ft" spoken by the agent and alerts me the same day

Risk	Guardrail
Existing-client complaint gets qualified like a lead (T09)	Detect "my project", "my designer" or a designer's name; transfer live in working hours; after hours, promise a senior callback by 10am and alert Nikhil
Callers hang up on a bot	Disclose up front, keep turns short, offer a person at any time; track hang-ups in the first 30 seconds
Misheard details (Marathi, Hinglish, noisy lines)	Read key facts back before booking ("a 3BHK in Kothrud, about 1,400 sq ft?"); test on real accents before launch
Wrong fit decision	Rules in code, not in the model; unclear cases go to the front-desk queue; weekly overturn rate on 10 random calls
Double-booked or unavailable designer	Free/busy checked at the moment of booking; unaccepted notes reassign after 30 min
Platform outage	Health check every 5 min; on failure, forwarding falls back to the front desk in hours and voicemail after hours
Caller data misuse	Spoken notice, paid AI tiers only, recordings behind login, a fixed retention period, deletion on request

Decisions I need from Nikhil before build

- ☐ Sign off qualified.md as the 8 written rules, including budget review
- ☐ Approve the price explanation, word for word
- ☐ Forward the existing studio number, or publish a new one?
- ☐ Which designer takes which areas and project types; when the principal designer goes on site
- ☐ A list of VIP referrers from his own network (T05)
- ☐ How long call recordings are kept
- ☐ Languages to support at launch: English, Hindi, Marathi

Research I still owe

- ☐ Vaani's all-in per-minute rate for Indian inbound calls, in writing
- ☐ September call records from the telco, for a phone-only baseline
- ☐ Telco charges and setup for conditional call forwarding on Aangan's plan

Extending to WhatsApp and the web form

Only the front door changes. The 8 rules, the Supabase record, the handoff note and the dashboard are built channel-neutral from day one, so each new channel is an adapter, not a rebuild.

Web form next (cheapest). The form already collects area, size, state, timeline and budget (F01–F10). The same extraction and rules run within a minute of submission; qualified enquiries get an email with live designer slots to pick from. "Please call me" (F04) goes to a human callback queue, and price-list requests (F07) get the same approved explanation in writing.

WhatsApp after that. Vaani lists WhatsApp as a channel for the same agent, so the script carries over. Aangan would need Meta business verification and approved message templates for replies outside the 24-hour window. Photos callers send (W09) attach to the designer's note. The September threads show the same after-hours gap as the phone: W01, W05 and W02 all waited a night, a weekend or two days.

Sources

Case material: Case 03 pre-read (Nikhil: Unanswered Enquiries), Aangan September 2026 enquiry transcripts, and the class slides "The Nine Checks" and "The Stack".

- [Lead response time: HBR and MIT/InsideSales findings](#) (TeleCRM)
- [Vaani Labs profile](#) (Voice AI Space)
- [Bolna pricing and India alternatives](#) (Edesy)
- [Bolna: buying Indian phone numbers](#) (Bolna docs)
- [Vapi and competitor pricing](#) (11x)
- [DPDP Act phased rollout and milestones](#) (AZB & Partners)
- [TRAI press release 11/2025 on TCCCPR amendments](#) (TRAI)
- [Gemini Developer API pricing](#) (Google)
- [Vercel fair use guidelines](#) (Vercel)

- [Cal.com pricing guide](#) (Zeeg)
- [Resend pricing](#) (Resend docs)